

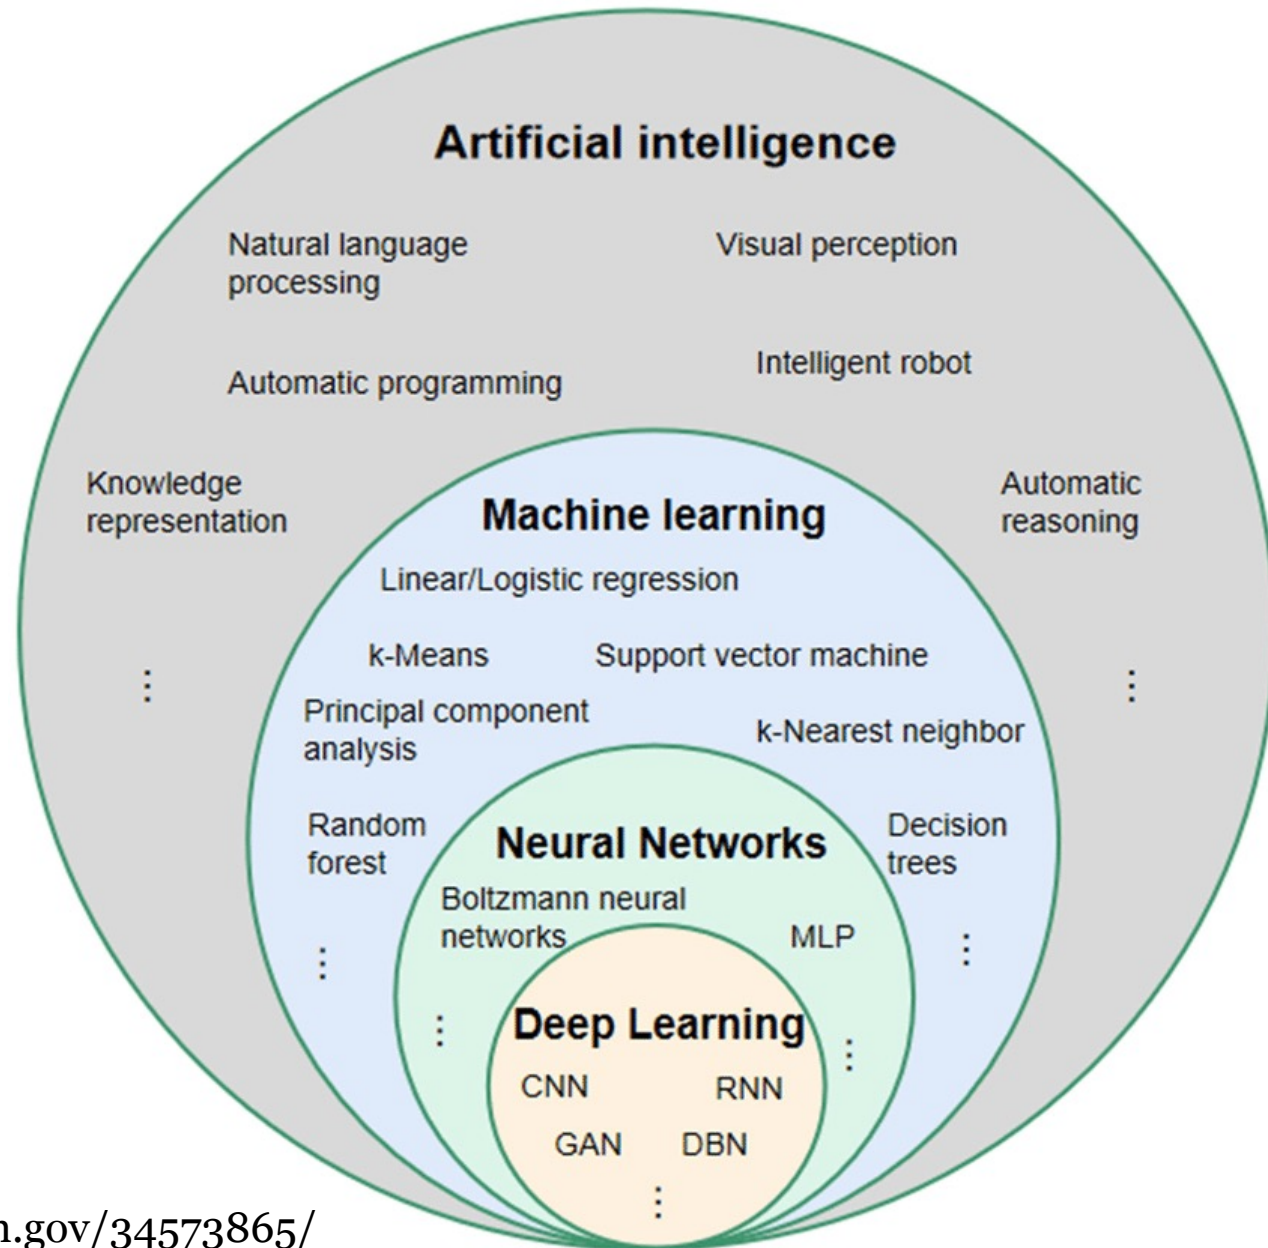
A Primer on Deep Learning

Dong Xu

*EECS Department
C. S. Bond Life Sciences Center
University of Missouri, Columbia
<http://digbio.missouri.edu>*

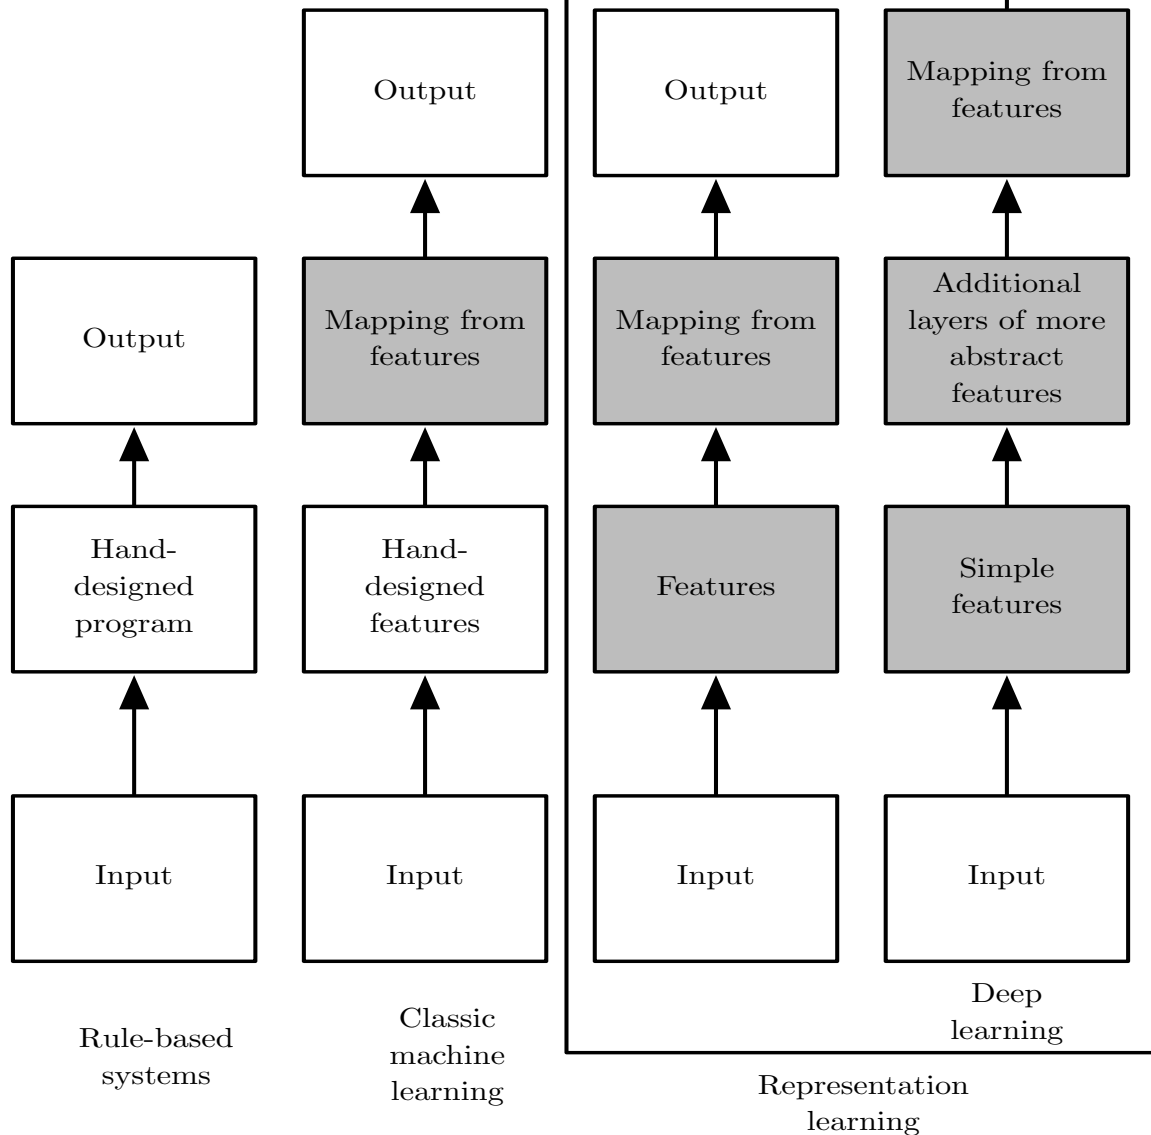


AI Scope



<https://pubmed.ncbi.nlm.nih.gov/34573865/>

Different Approaches



Supervised Machine learning

- Apply a prediction function to a feature representation of the image to get the desired output:


$f(\text{apple image}) = \text{"apple"}$

$f(\text{tomato image}) = \text{"tomato"}$

$f(\text{cow image}) = \text{"cow"}$

Supervised Machine learning

$$y = f(\mathbf{x})$$



- **Training:** given a *training set* of labeled examples $\{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$, estimate the prediction function f by minimizing the prediction error on the training set
- **Testing:** apply f to a never before seen *test example* \mathbf{x} and output the predicted value $y = f(\mathbf{x})$

Generalization



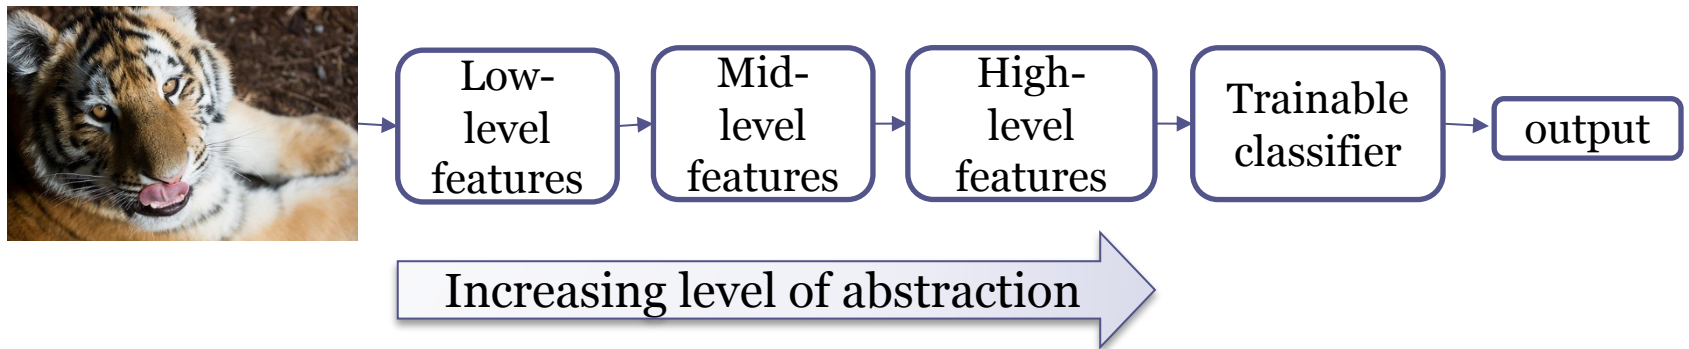
Training set (labels known)



Test set (labels unknown)

- How well does a learned model generalize from the data it was trained on to a new test set?


Learning Hierarchical Representations



- Hierarchy of representations with increasing level of abstraction. Each stage is a kind of trainable nonlinear feature transform
- Image recognition
Pixel → edge → texton → motif → part → object
- Text
Character → word → word group → clause → sentence
→ story


Example Application

- Handwriting Digit Recognition



Handwriting Digit Recognition


Input



Ink \rightarrow 1


No ink \rightarrow 0

Output



Each dimension represents the confidence of a digit.

Deep Neural Network (DNN)



Deep means many hidden layers

Abstraction and Representation

physical space \rightarrow latent/embedding space (manifold)



photo



impressionism



expressionism





cubism




abstract
expressionism

Latent Representation



Neuron

`layern[x,y,z]`



Channel

`layern[:, :, z]`




Layer/DeepDream

`layern[:, :, :]2`



Class Logits

`pre_softmax[k]`




Class Probability

`softmax[k]`

Unsupervised Learning & Embedding


High-dimensional signals
in a signal space S



Mapping
(by learning in DL)

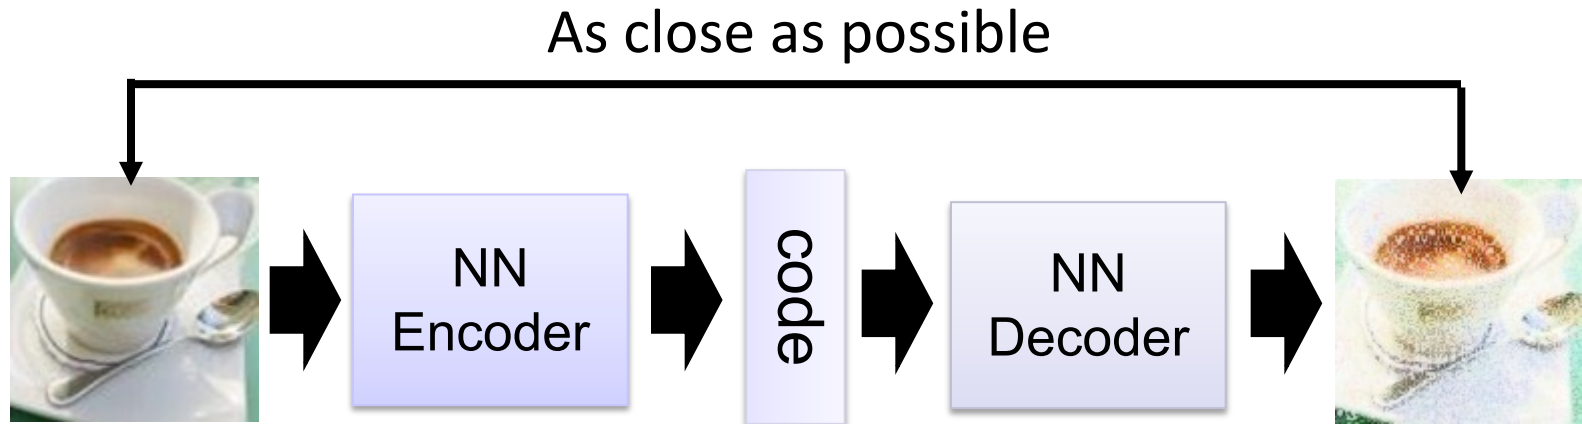


Embedding space
(e.g., low-dimension,
small number of bits)




An embedding is a function from an original space to an embedding space that preserves aspects of the geometry of the original space


Autoencoder



A type of unsupervised learning which discovers generic features of the data (**learn data patterns**)



Autoencoder with 3 Layers

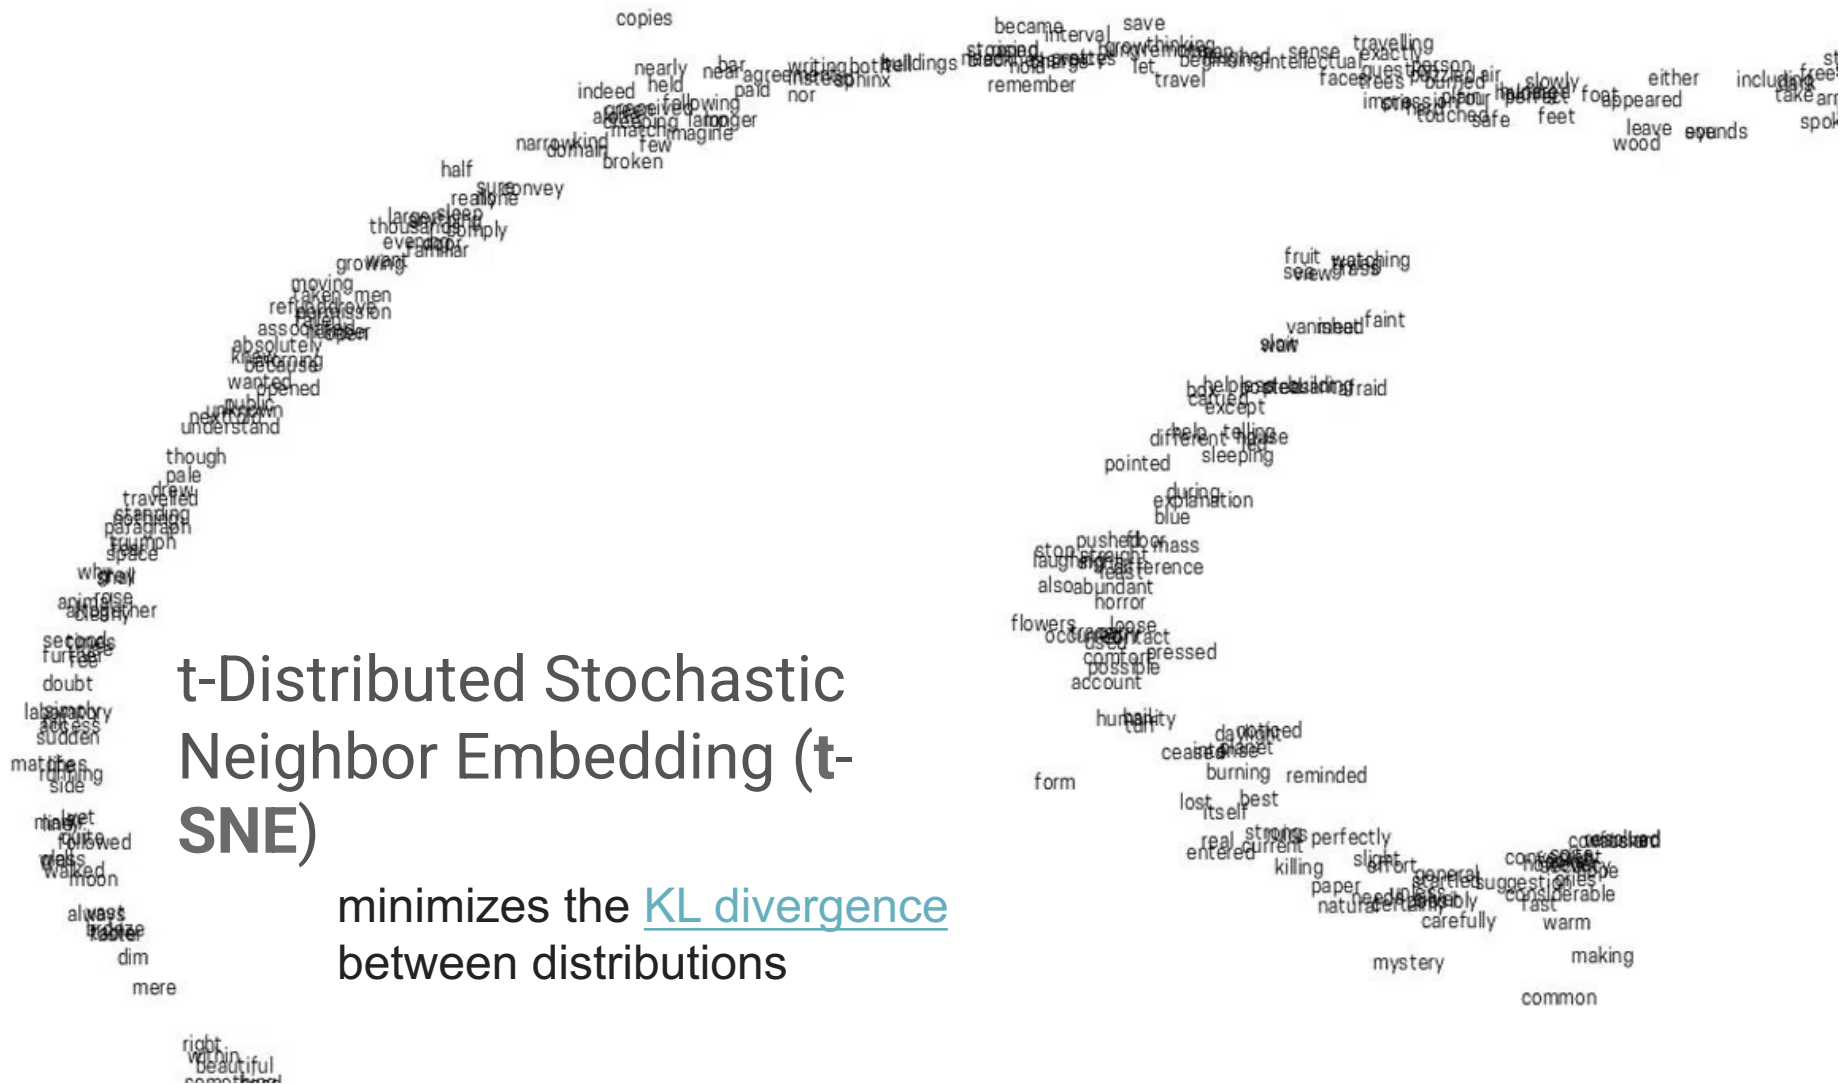


Large \rightarrow small, learn to compress

Z: Latent variable; embedding

Training: $\text{model.fit}(X, X)$
Cost function: $\sum_{k=1..N} (x_k - x'_k)^2$

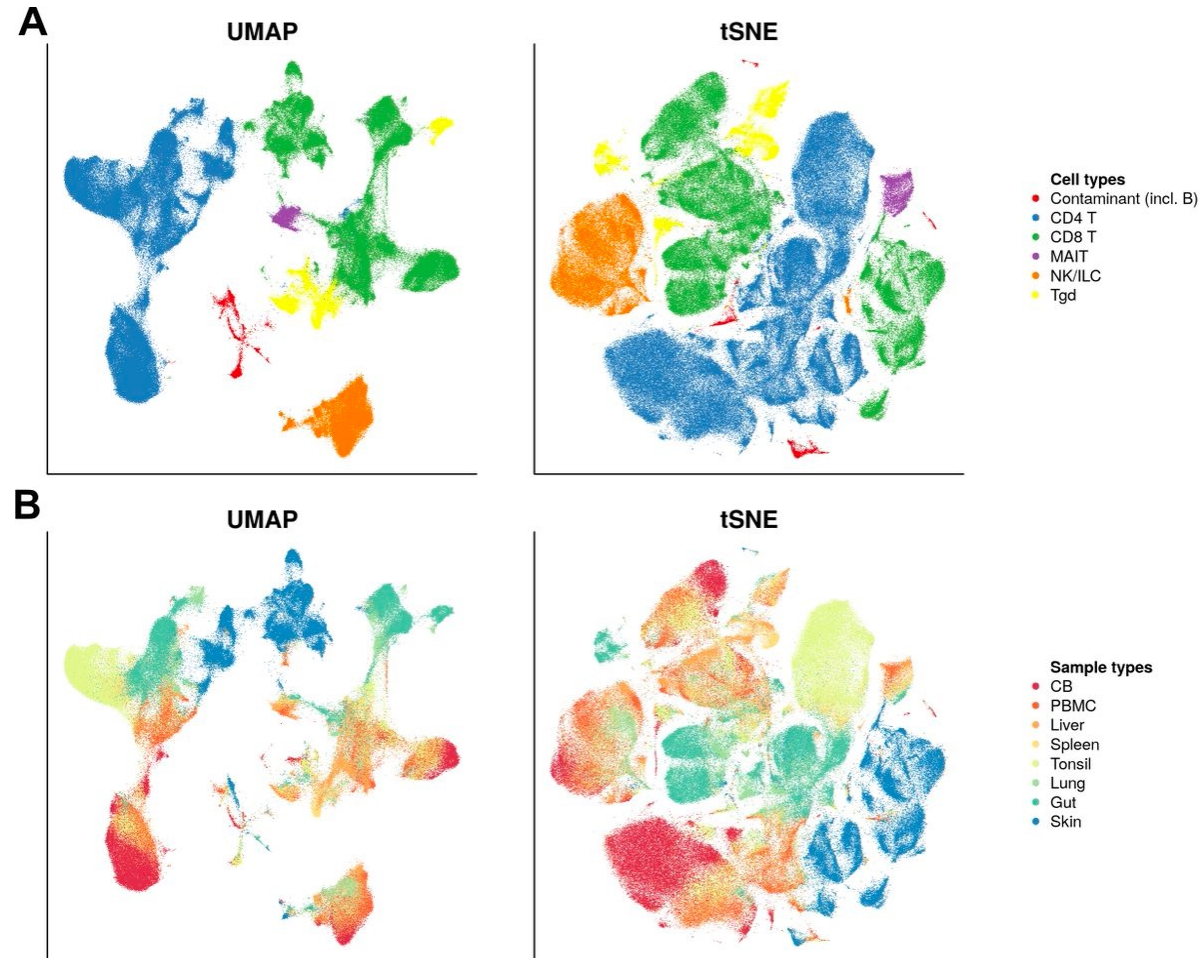
T-SNE Visualization




U-Map

Uniform Manifold Approximation and Projection (UMAP)

find a topological representation of the data in a lower dimensional space through manifold learning technique




Manifold




- A Manifold is a topological space that locally resembles Euclidean space near each point
- n-dim manifold \rightarrow topological space M , every point $x \in M$ has a neighbor homeomorphic (isomorphic) to Euclidean space \mathbb{R}^n

Manifold Hypothesis

- **DL Central Hypothesis:** Data concentrates around a low-dim manifold (**relevant dimension**)
- Mimic human learning
- Not all embeddings produce manifold



Self-supervised learning



representation learning that enables learning good data representation from unlabeled dataset.

Foundation model era

- Machine learning paradigms

Feature engineering: manual feature extraction (SVM, LightGBM, XGBoost)

Architecture engineering: raw features, design deep network (CNN, LSTM)

Objective engineering: pre-train large model and fine-tune it (ResNet50, Bert)

Prompt engineering: prompt **foundation model** in zero/few shots

- Industrial era of artificial intelligence

